

Azure Cosmos DB

The Swiss Army NoSQL Cloud Database

codit|



Thank you to our sponsors!

Gold Sponsors



Silver Sponsors



Community Sponsors



What can you expect?

- Introduction of Cosmos DB
- Scenarios
- Cosmos DB Key Characteristics
- Multi-model
- Reference case
- Demos



Cosmos DB

Introduction

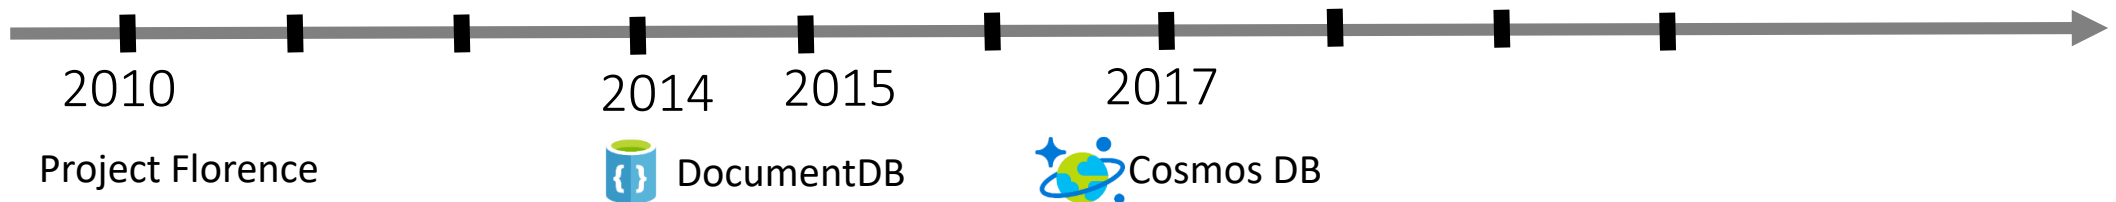
Scenarios



Azure Cosmos DB Evolution



- Originally started to address the problems faced by large scale apps inside Microsoft
- Built from the ground up for the cloud
- Used extensively inside Microsoft
- One of the fastest growing services on Azure





Cosmos DB Service

A globally distributed, massively scalable, multi-model database service



Table API



MongoDB API



cassandra
Cassandra API



Key-value



Column-family



Document



Graph

Elastic scale out
of storage & throughput

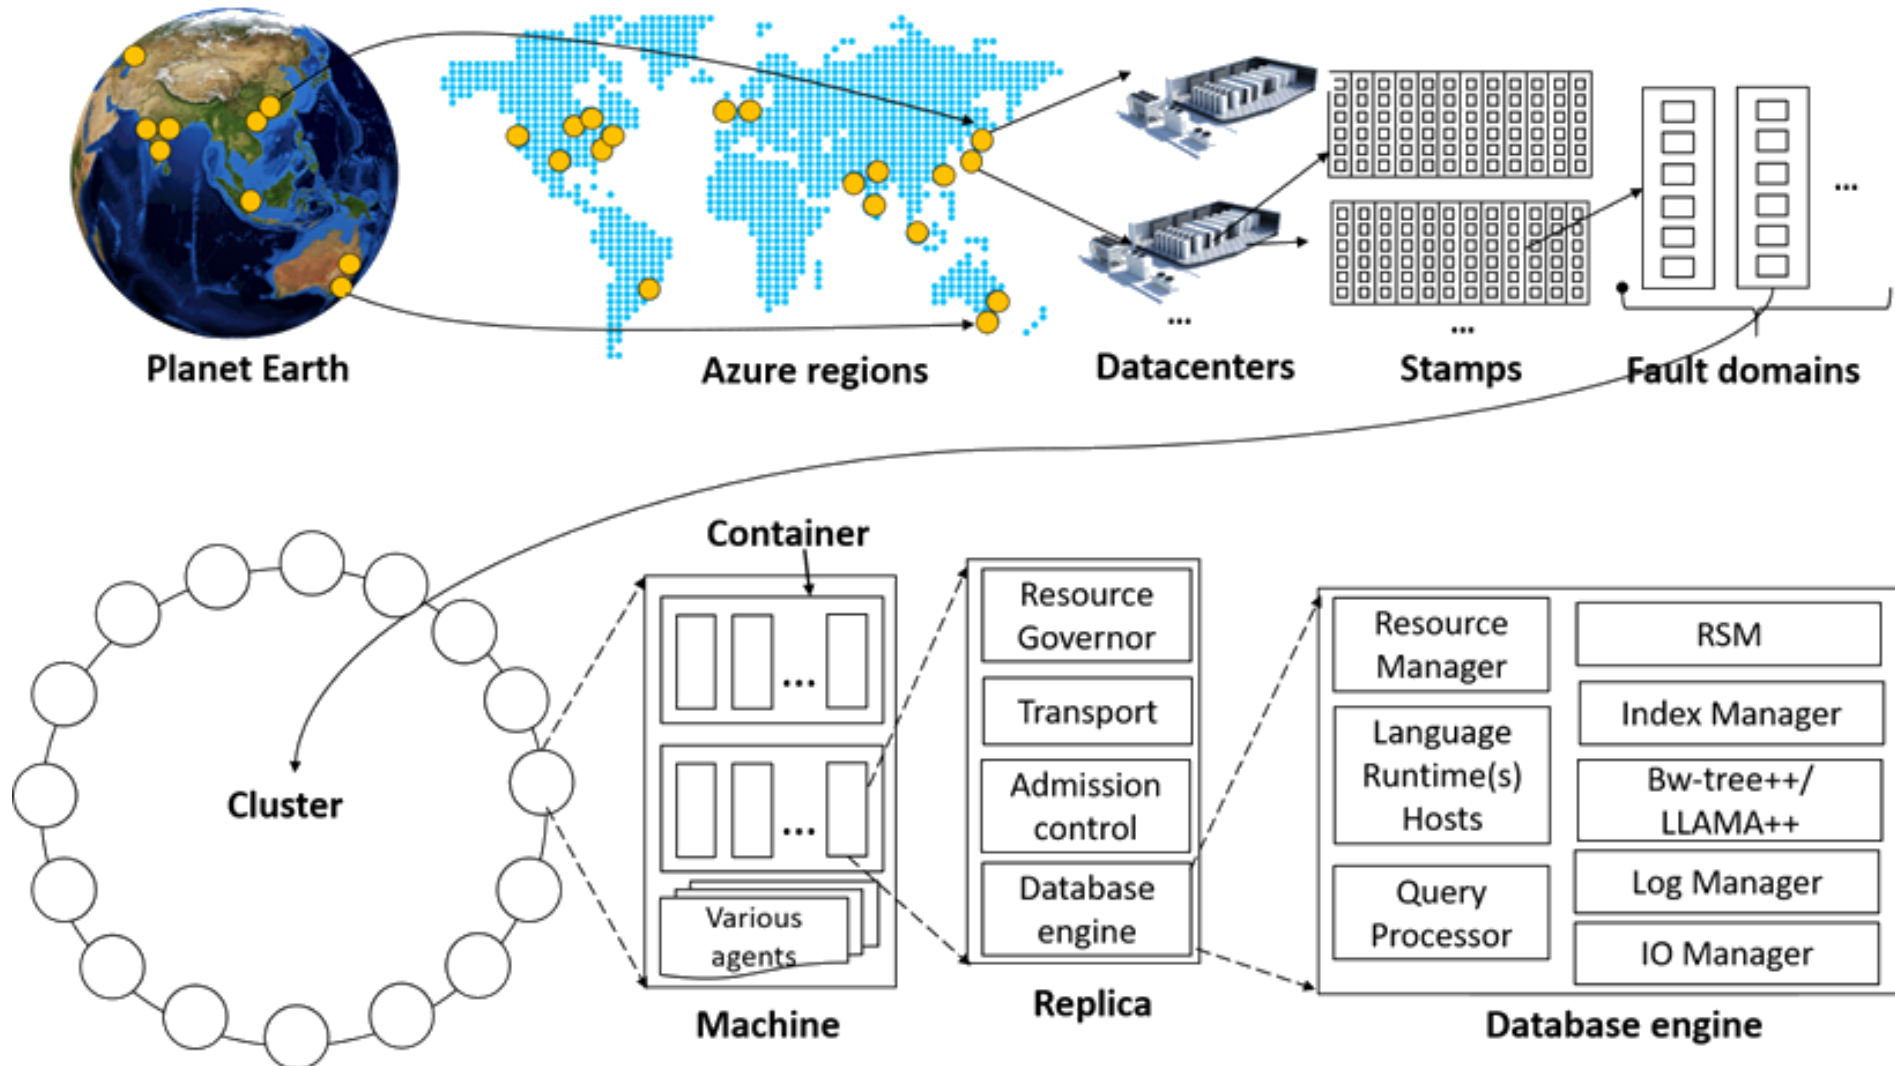
Guaranteed low latency at the 99th percentile

Five well-defined consistency models

Turnkey global distribution

Comprehensive SLAs

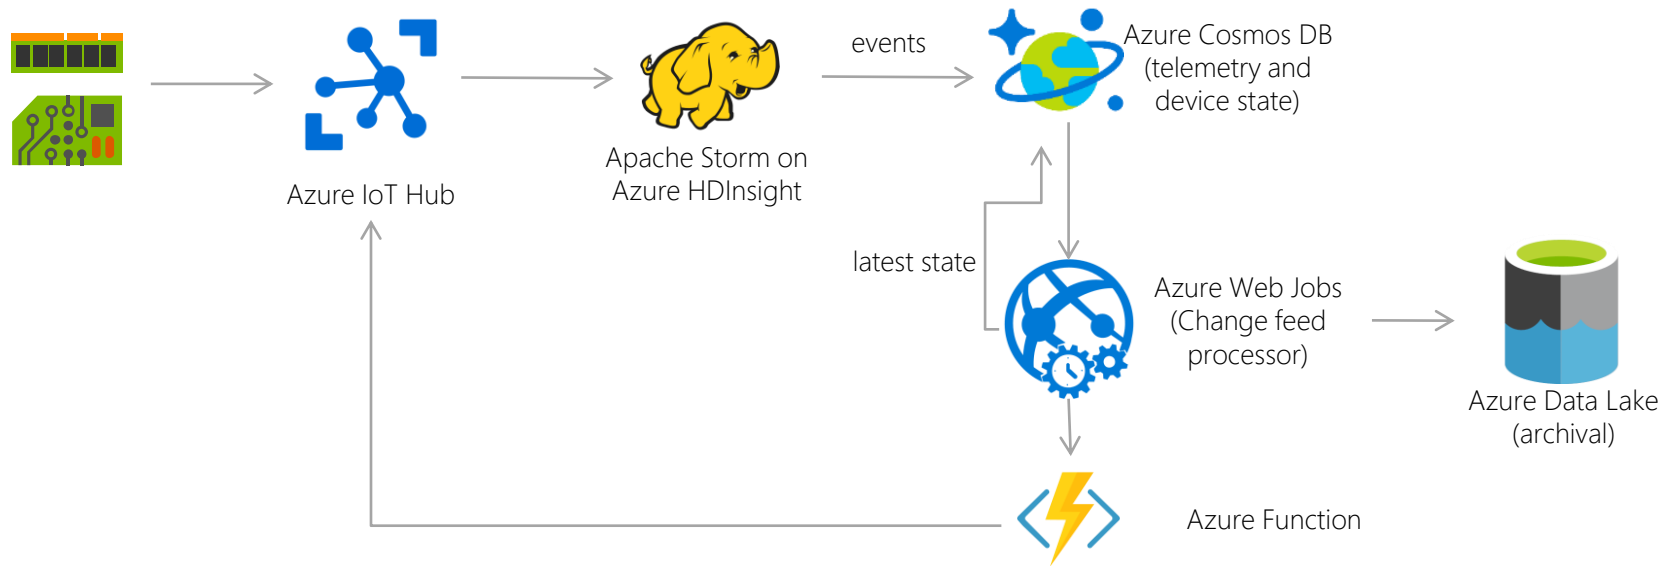
System topology (behind the scenes)



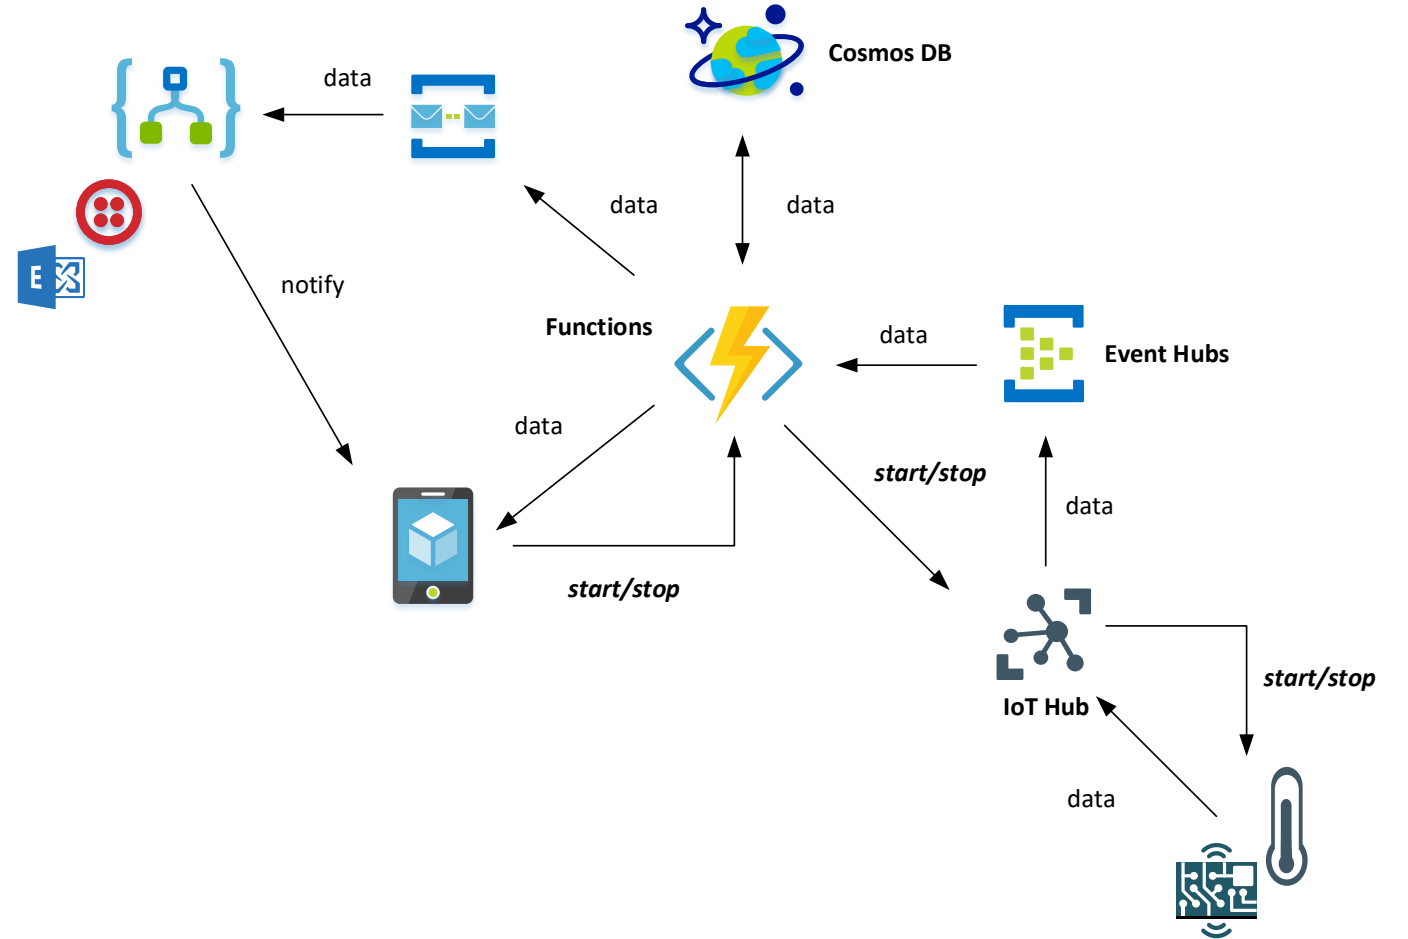
Customers



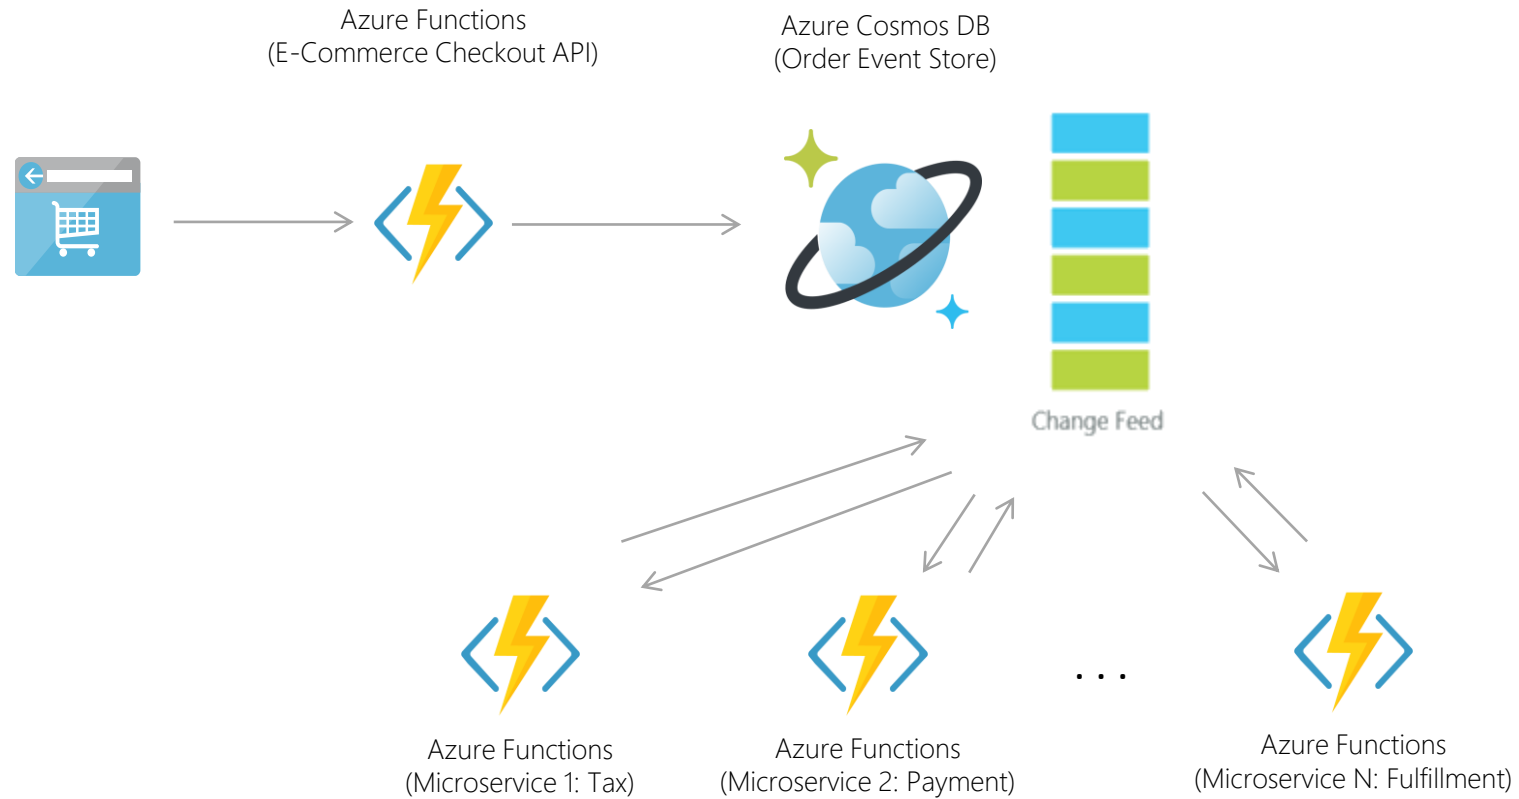
Scenario - Telemetry & Sensor Data



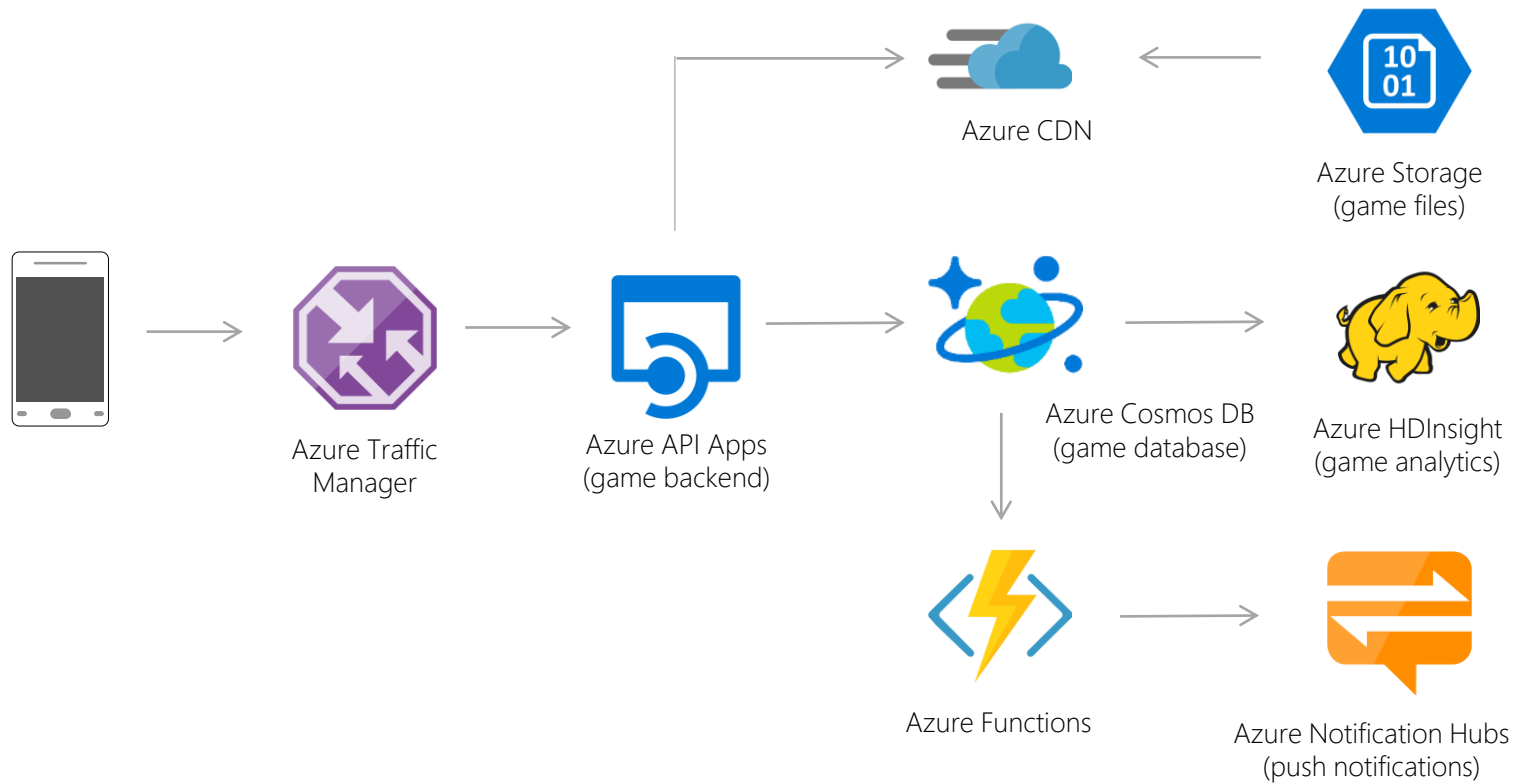
Demo – IoT



Scenario – Order processing



Scenario – Multiplayer gaming



Cosmos DB Characteristics

Global distribution

Resource model

Scale

Consistency

Indexing



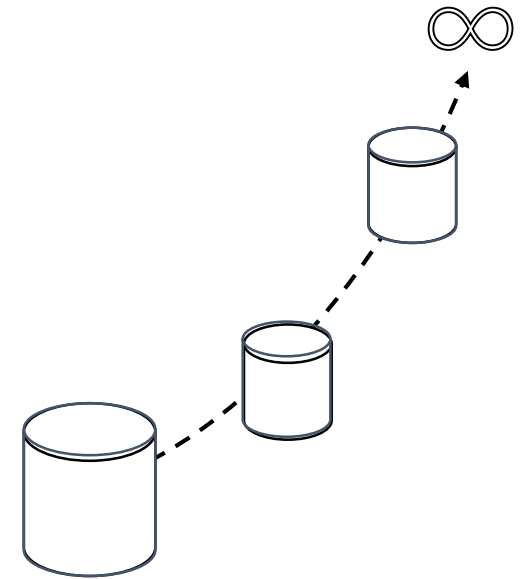
Global Distribution

High Availability

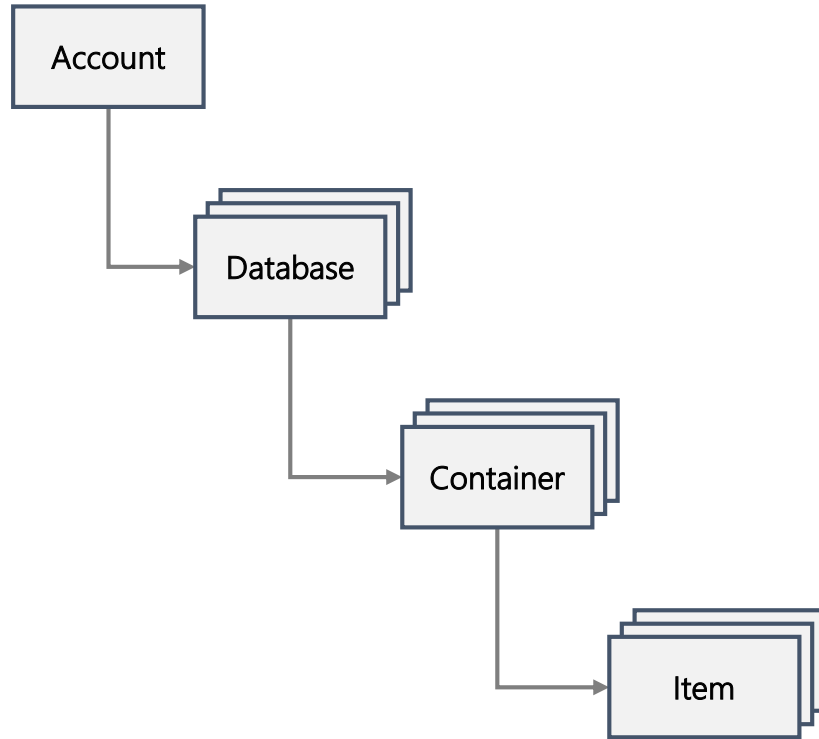
- Automatic and Manual Failover
- Multi-homing API removes need for app redeployment

Low Latency (anywhere in the world)

- Sending a packet across the world under ideal network conditions takes 100's of milliseconds
- Packets cannot move faster than the speed of light

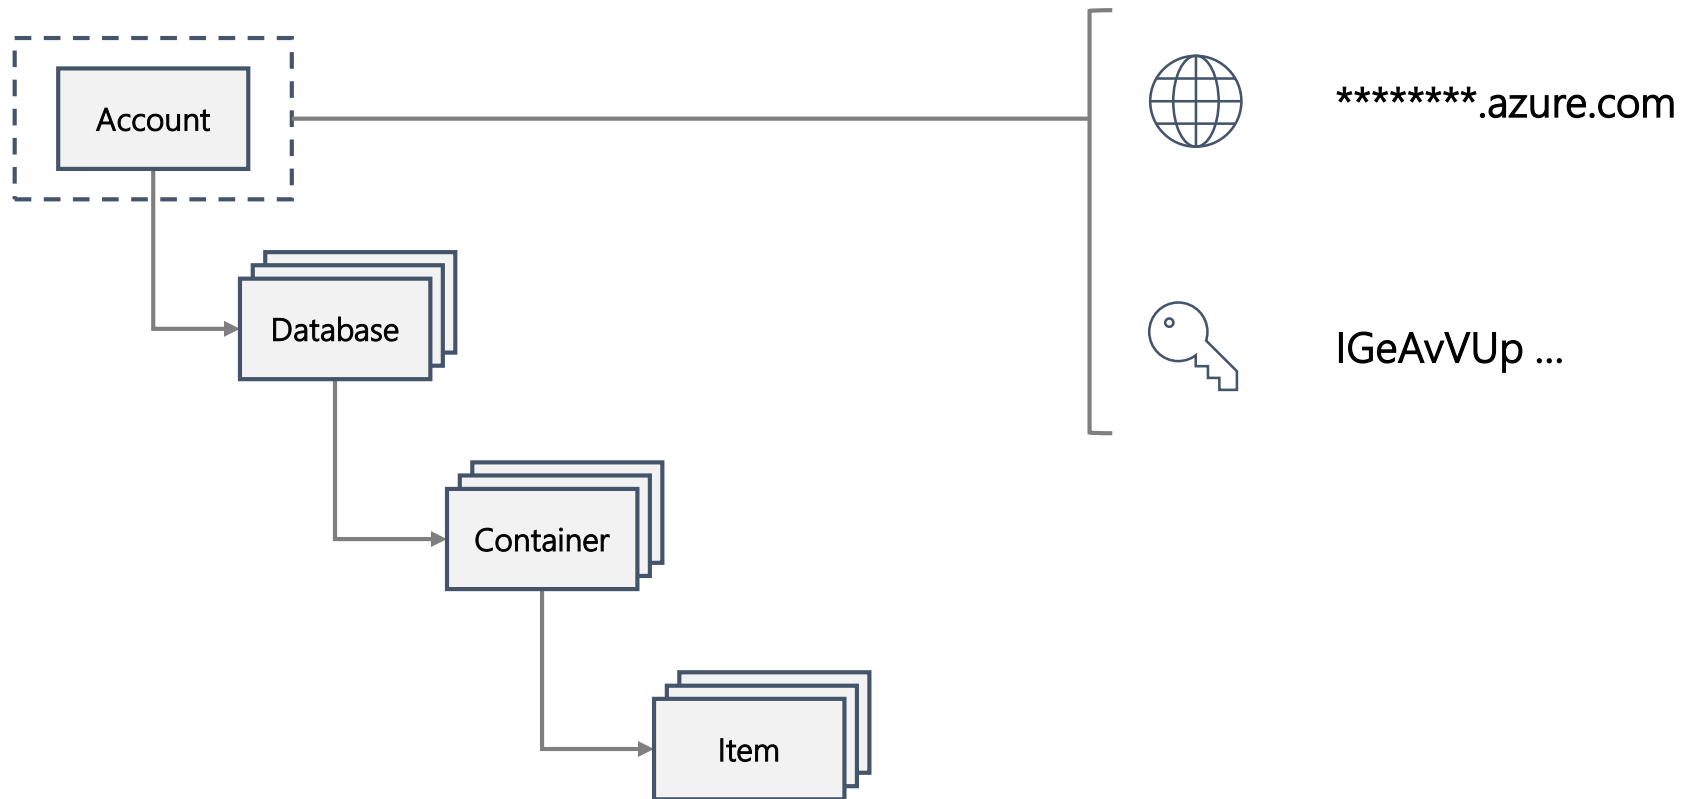


Resource Model

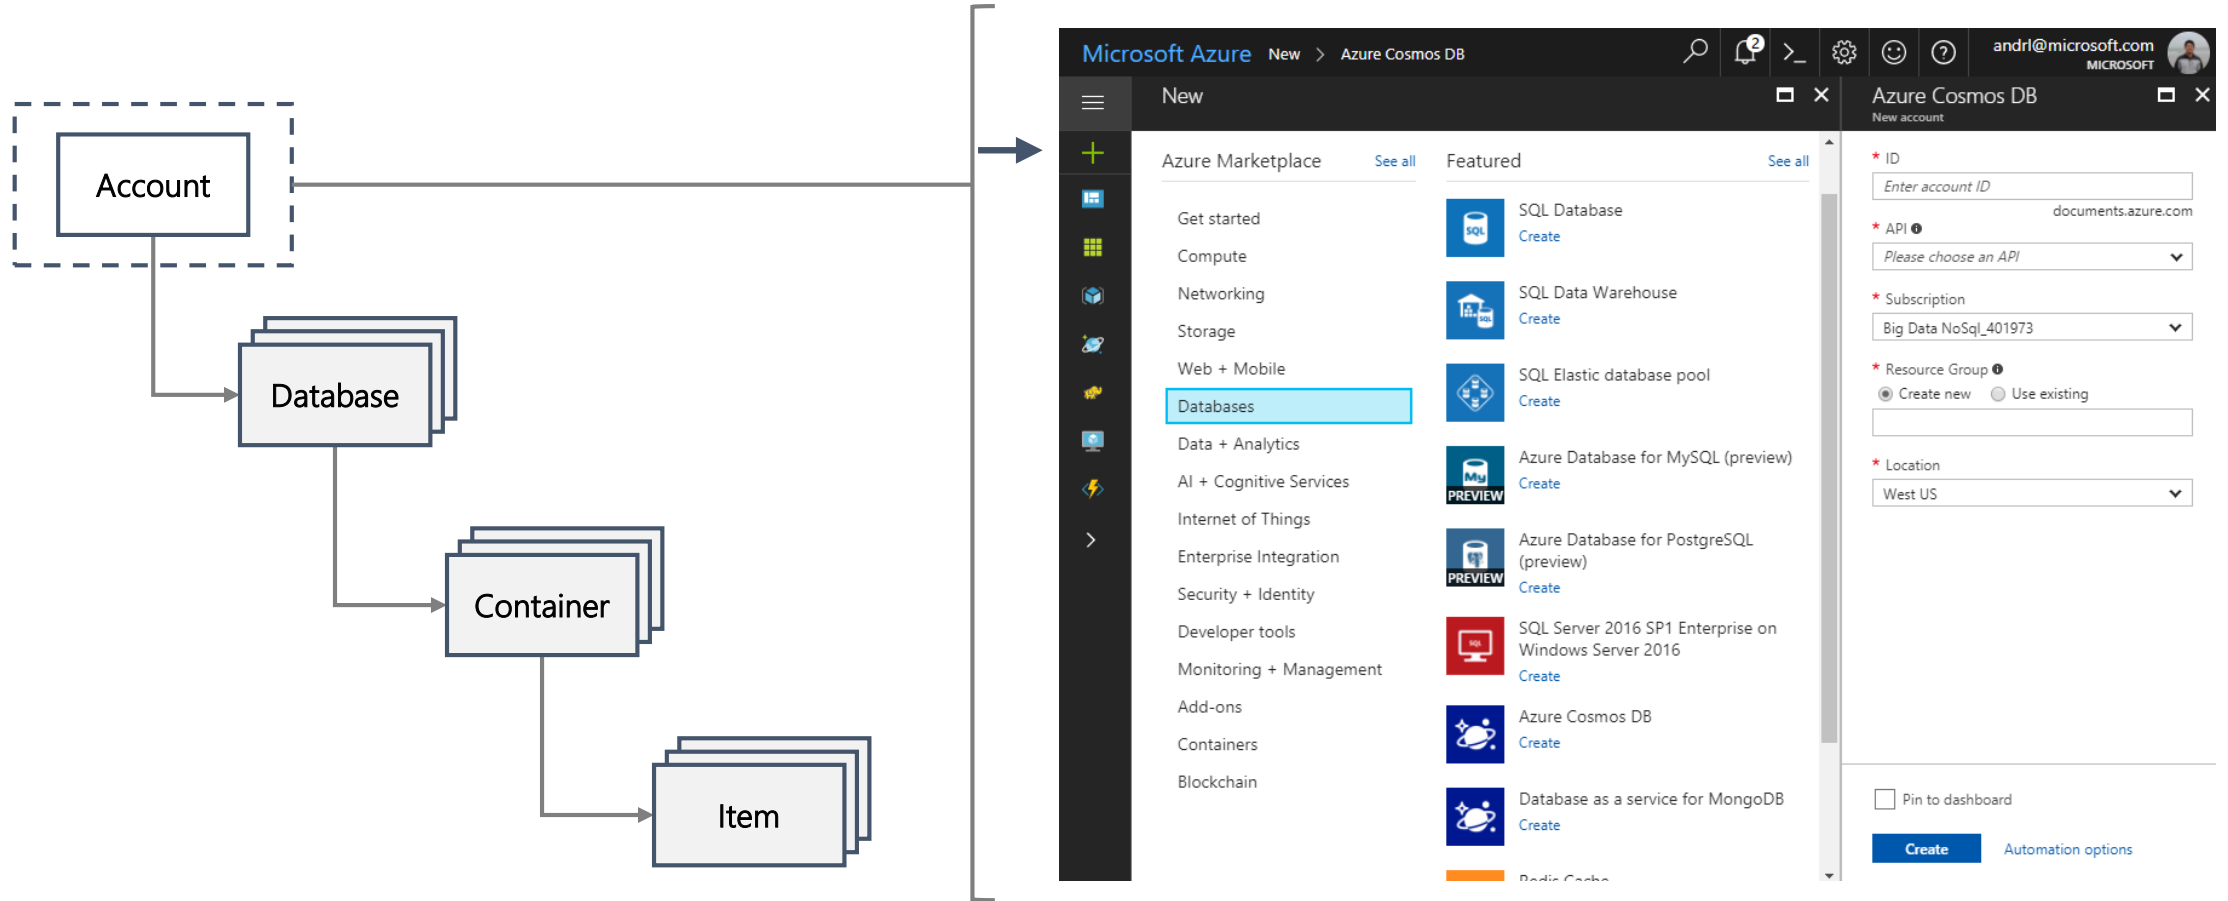


- Resources identified by their logical and stable URI
- Hierarchical overlay over horizontally partitioned entities; spanning machines, clusters and regions
- Extensible custom projections based on specific type of API interface
- Stateless interaction (HTTP and TCP)

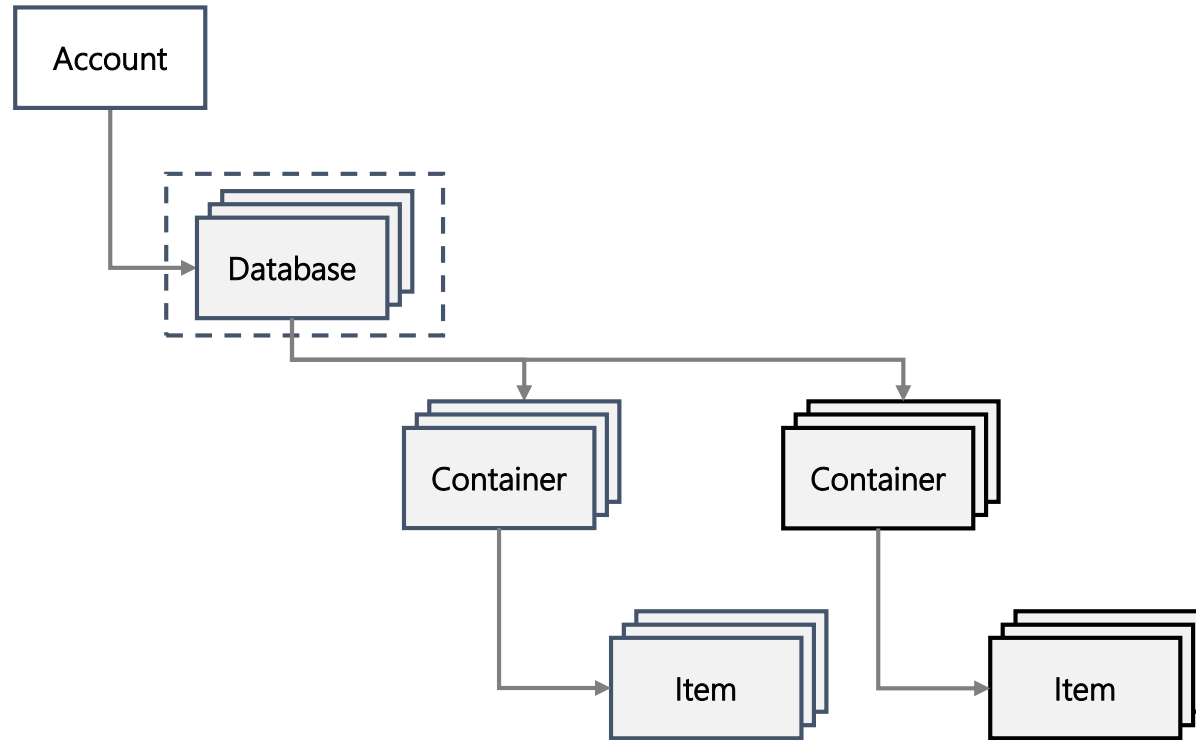
Account URI and Credentials



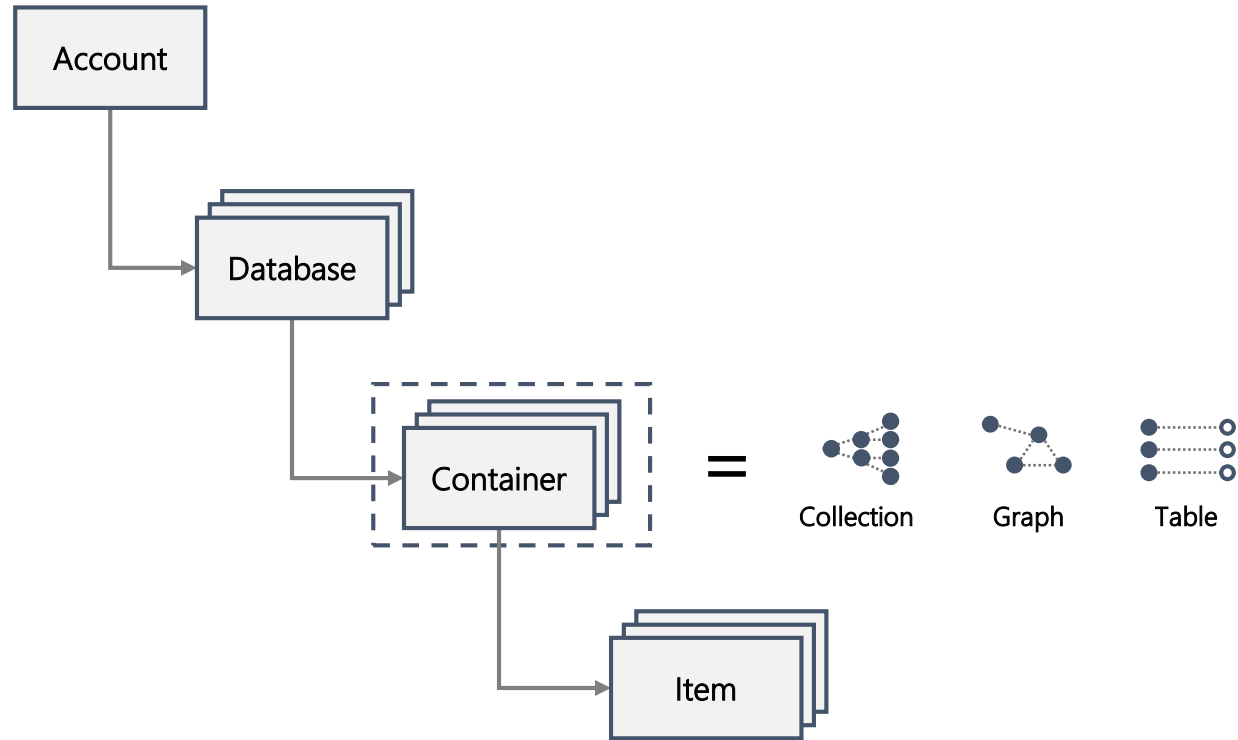
Creating Account



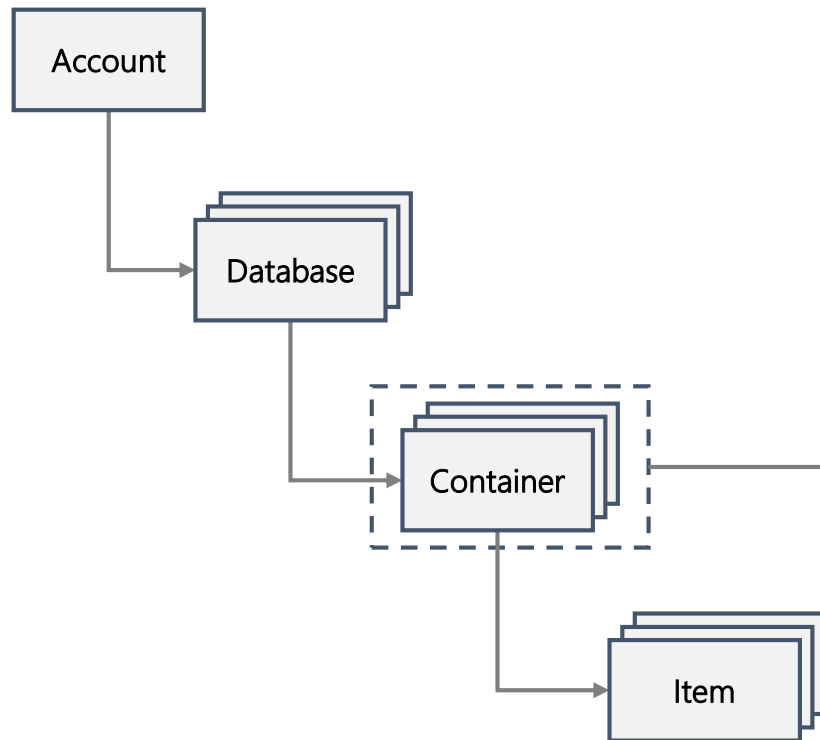
Database representation



Container representation



Creating collection – SQL API



The screenshot shows the 'Add Collection' dialog box in the Azure Cosmos DB SQL API. The dialog has a title bar with 'Add Collection' and window control buttons. The main content area contains several fields and options:

- Collection Id**: A text input field with the placeholder 'Enter collection id'. A blue arrow points to this field.
- STORAGE CAPACITY**: Two radio buttons, 'Fixed (10GB)' and 'Unlimited*'. The 'Unlimited*' option is selected. Below the buttons, a note states: '*up to 10TB, request higher capacity via support.'
- INITIAL THROUGHPUT CAPACITY (RU/s)**: A numeric input field with the value '10000' and a green checkmark. To the right are minus and plus buttons. Below the field, a note states: 'Between 2500 and 100000. You can provision higher throughput capacity via support request. [Click here.](#) Estimated hourly spend \$0.800USD'.
- PARTITION KEY**: A text input field with the value '/deviceid' and a green checkmark.
- DATABASE**: Two radio buttons, 'Create New' and 'Use existing'. The 'Use existing' option is selected. Below the buttons is a dropdown menu with the value 'andri-dev'.

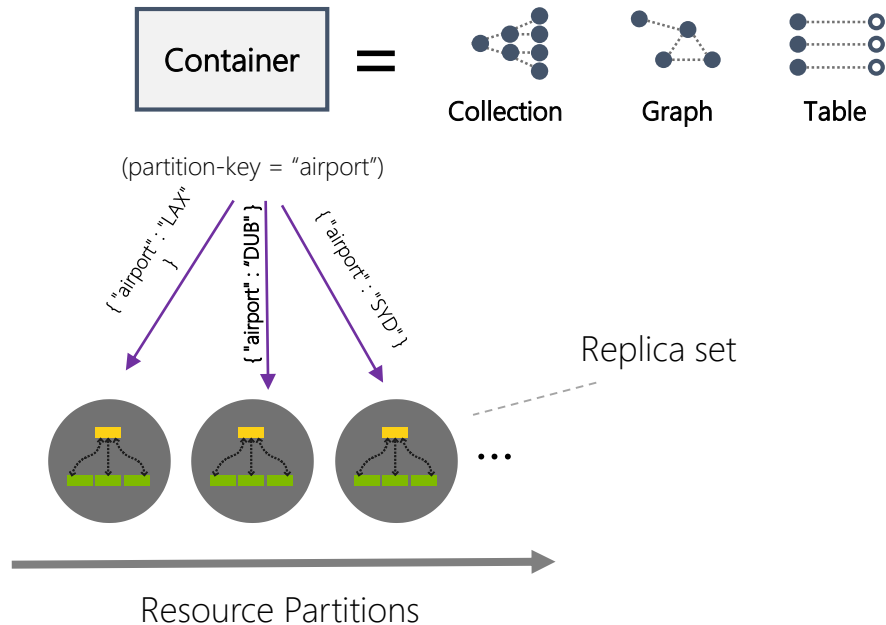
At the bottom of the dialog is a blue 'OK' button.

Scale

- Pay as go for storage and throughput
- Elastic scale across regions
- Partitions



Horizontal Scaling

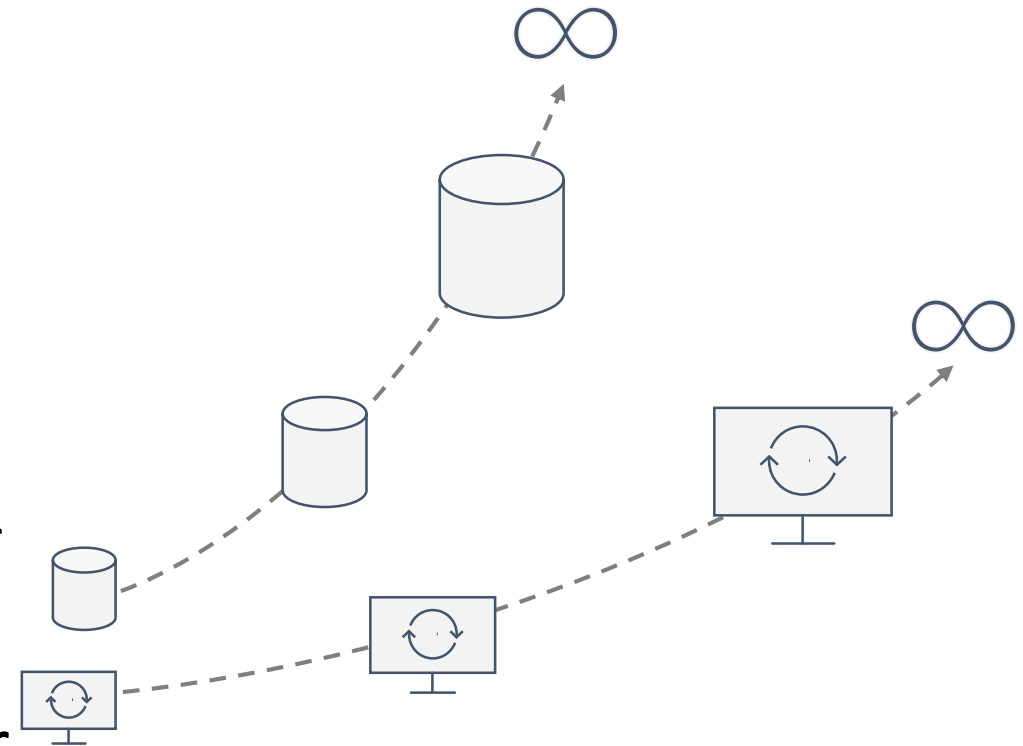


- Containers are horizontally partitioned
- Each partition made highly available via a replica set
- Partition management is transparent and highly responsive
- Partitioning scheme is dictated by a “partition-key”

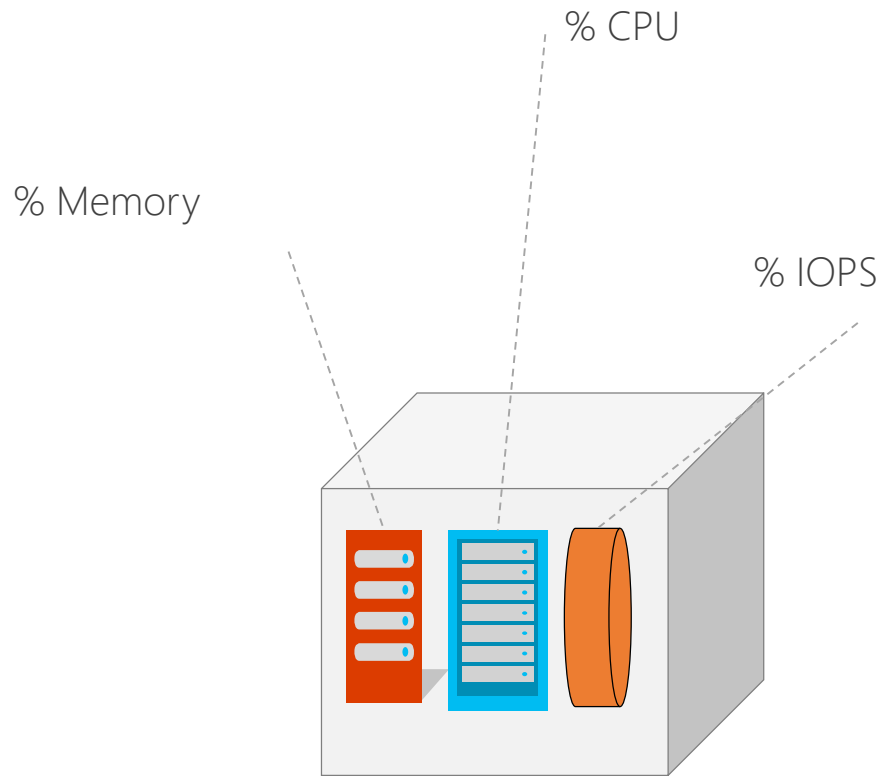
Elastic Scale Out of Storage and Throughput

SCALES AS YOUR APPS' NEEDS CHANGE

- Database elastically scales storage and throughput
- How? Scale-out!
- Collections can span across large clusters of machines
- Can start small and seamlessly grow as your app grows

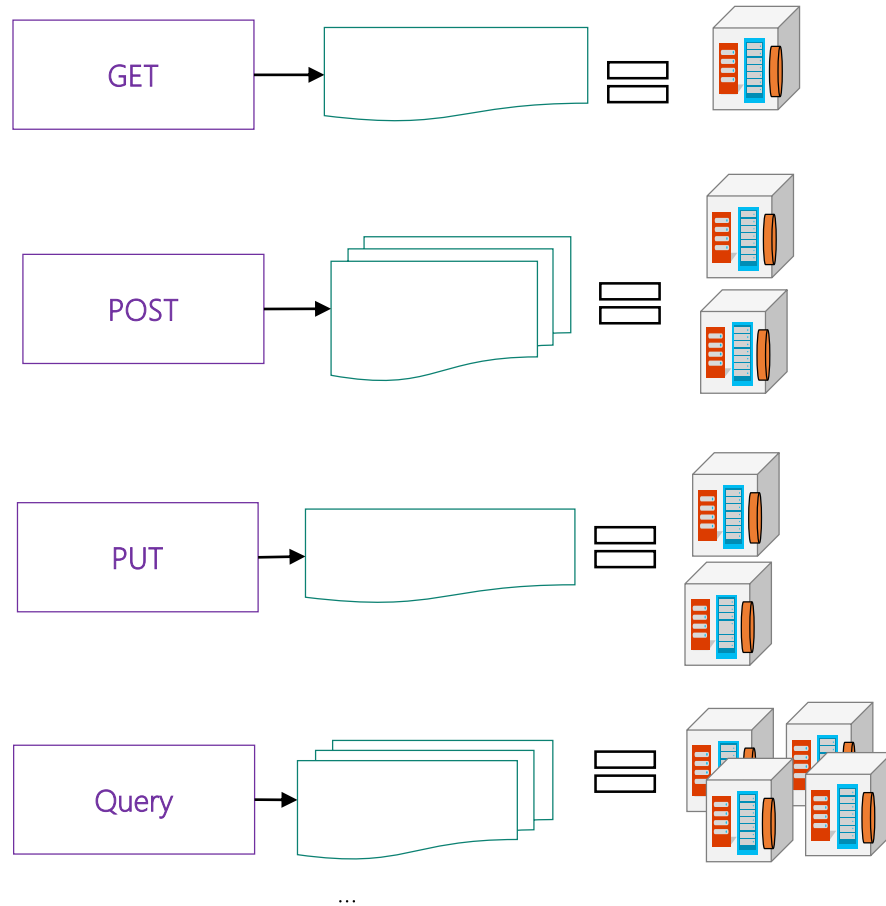


Request Units



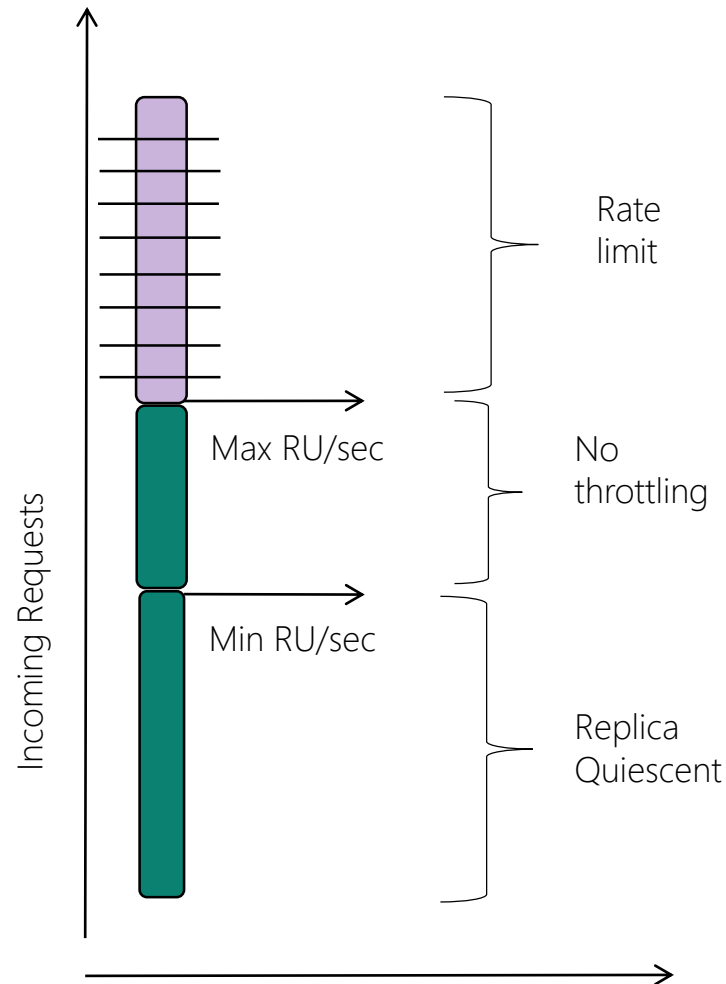
- Request Units (RU) is a rate-based currency
- Abstracts physical resources for performing requests
- Key to multi-tenancy, SLAs, and COGS efficiency
- Foreground and background activities

Request Units



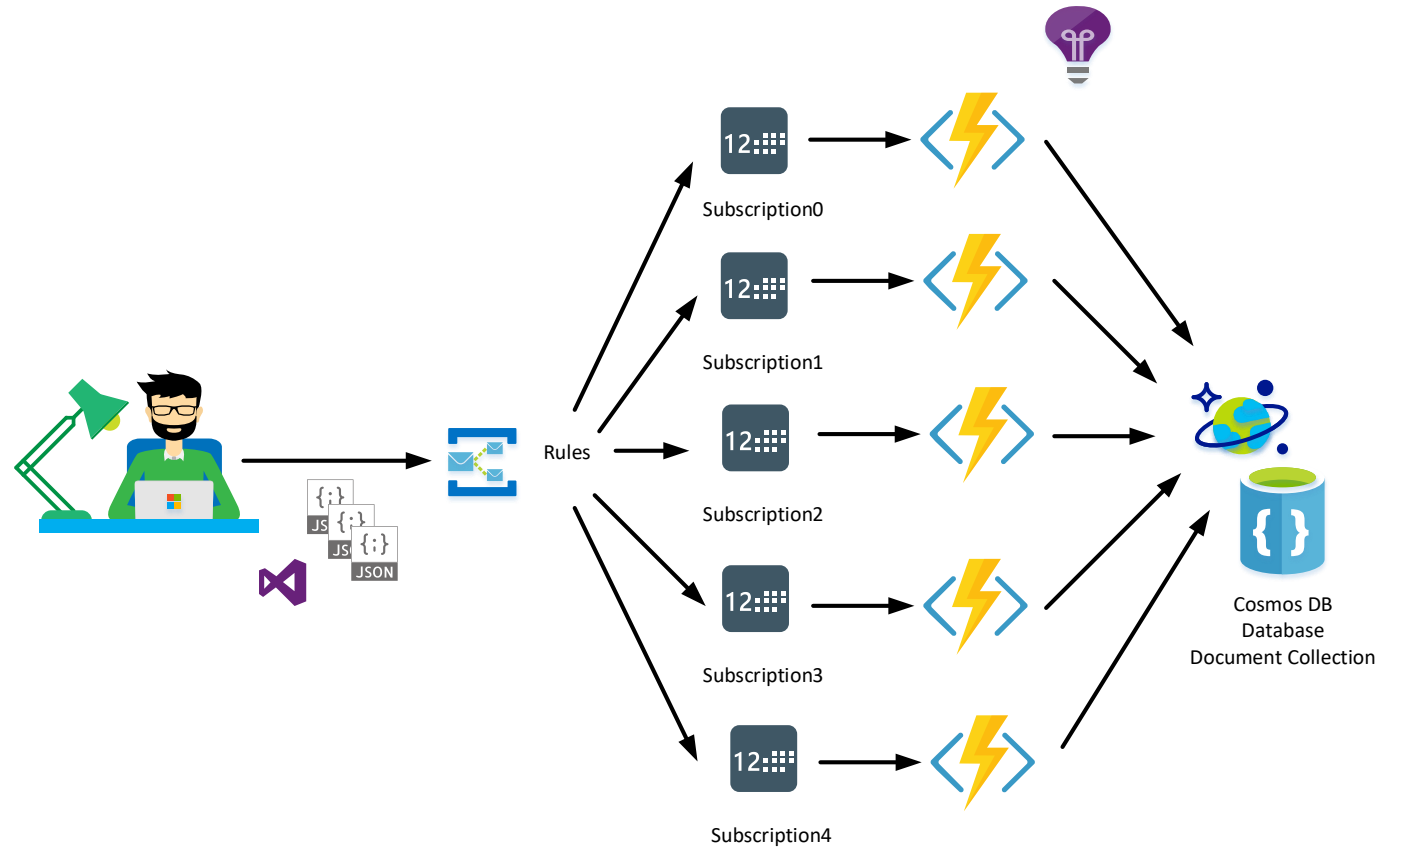
- Normalized across various access methods
- 1 RU = 1 read of 1 KB document
- Each request consumes fixed RUs
- Applies to reads, writes, queries, and stored procedure execution

Request Units

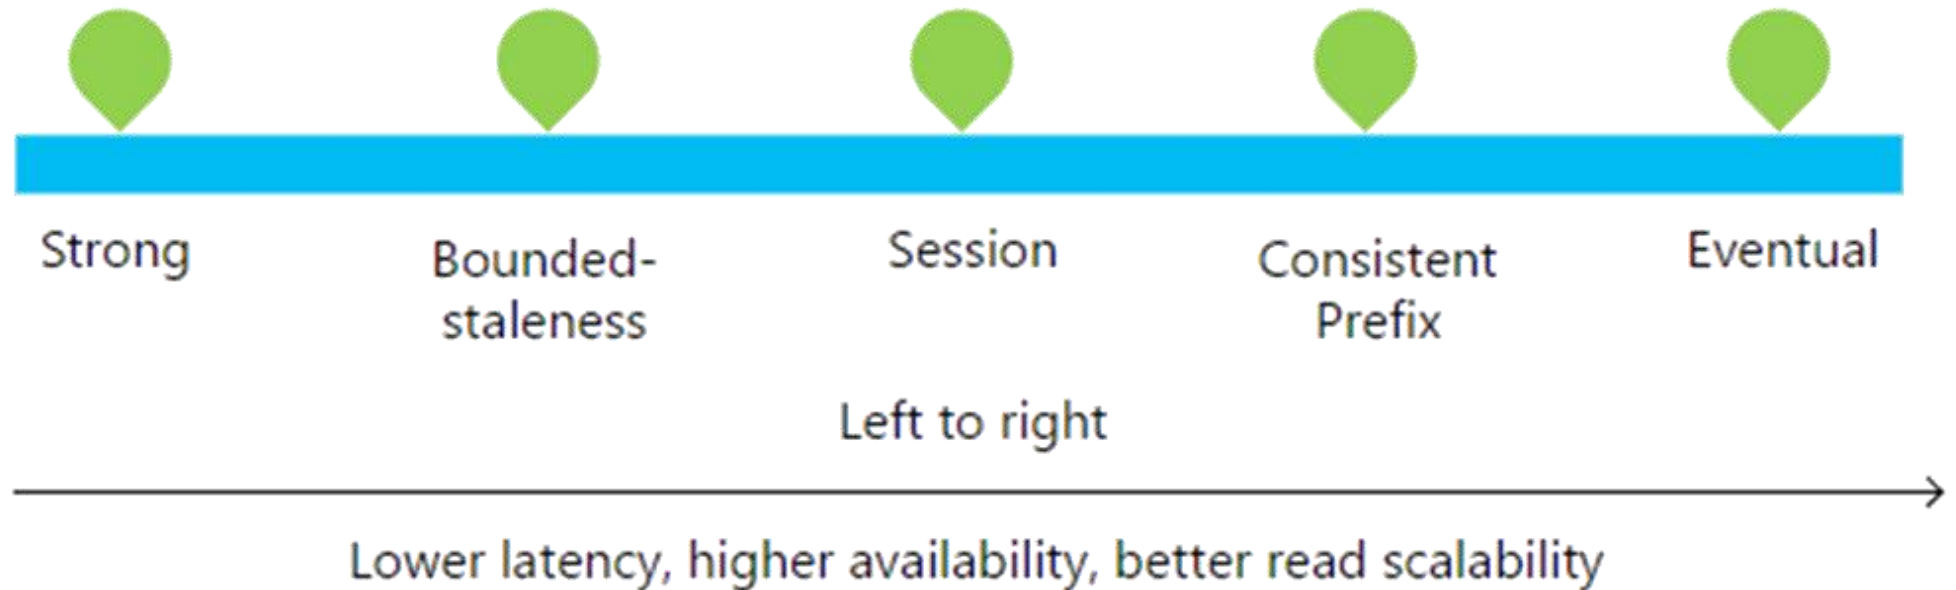


- Provisioned in terms of RU/sec and RU/min granularities
- Rate limiting based on amount of throughput provisioned
- Can be increased or decreased instantaneously
- Metered Hourly
- Background processes like TTL expiration, index transformations scheduled when quiescent

Demo - Scale



5 Well-defined, consistency models



Choose the right consistency

- Overridable on a per-request basis
- Provides control over performance-consistency tradeoffs, backed by comprehensive SLAs.
- An intuitive programming model offering low latency and high availability for your planet-scale app.

CLEAR TRADEOFFS

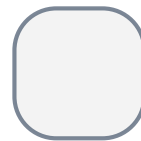
- Latency
- Availability
- Throughput



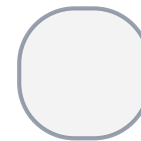
Strong



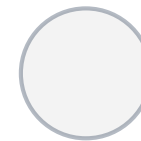
Bounded-staleness



Session



Consistent prefix



Eventual



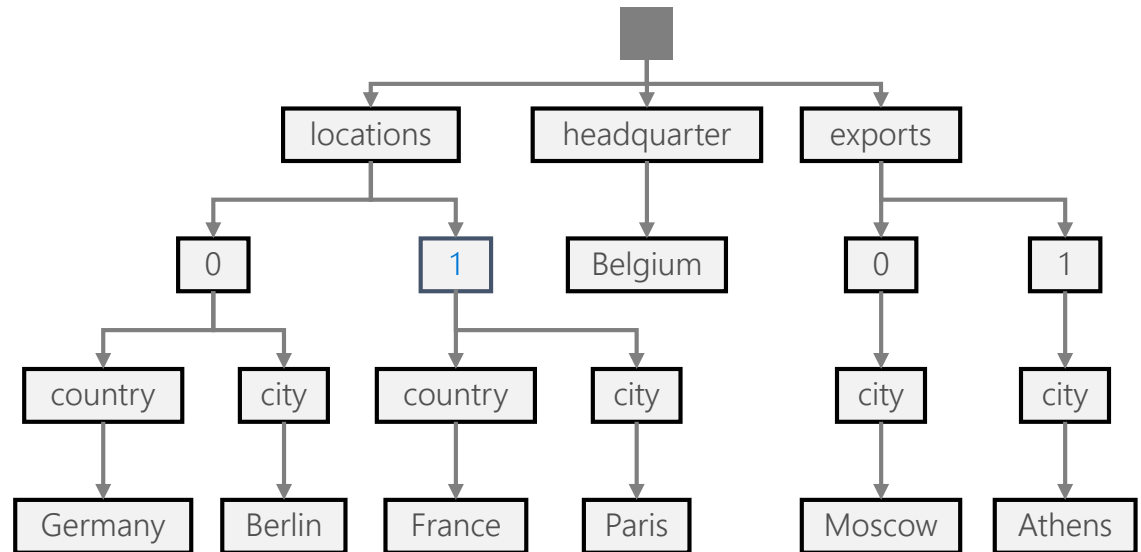
Handle any data with no schema or indexing required

Azure Cosmos DB's schema-less service automatically indexes all your data, regardless of the data model, to deliver fast queries.

- Automatic index management
- Synchronous auto-indexing
- No schemas or secondary indices needed
- Works across every data model

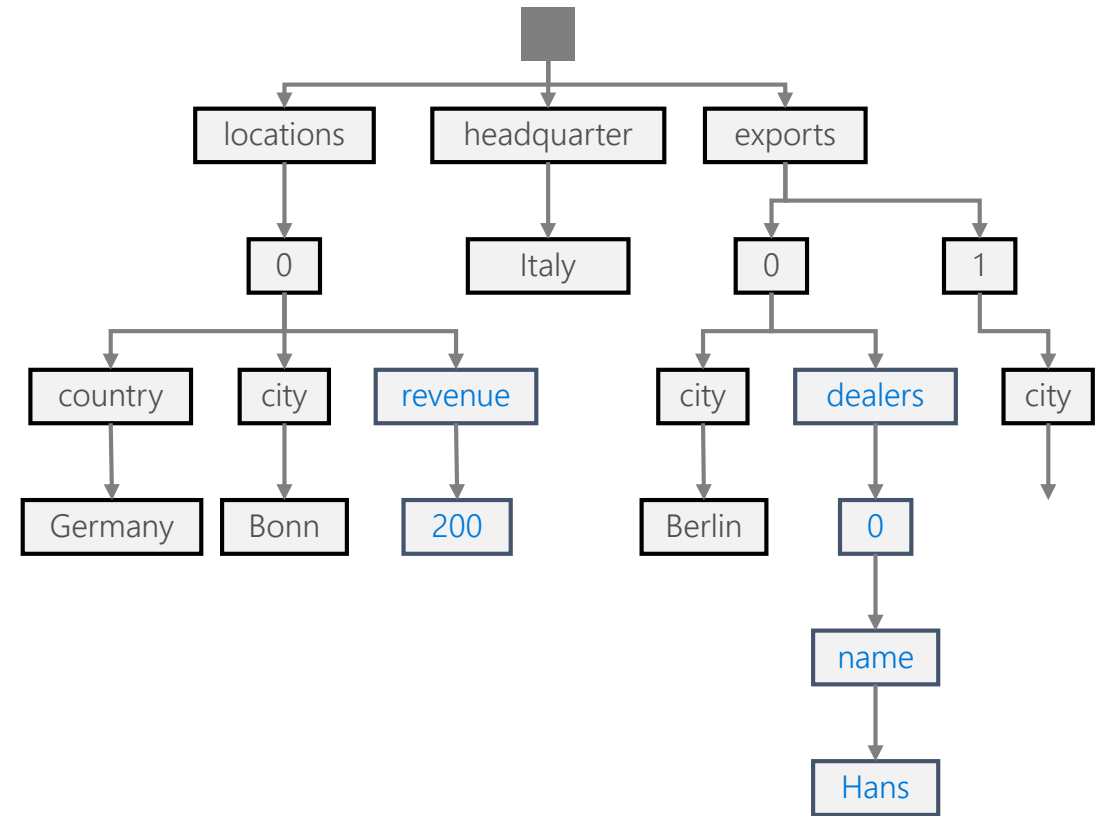
Indexing Json Documents

```
{  
  "locations": [  
    {  
      "country": "Germany",  
      "city": "Berlin"  
    },  
    {  
      "country": "France",  
      "city": "Paris"  
    }  
  ],  
  "headquarter": "Belgium",  
  "exports": [  
    { "city": "Moscow" },  
    { "city": "Athens" }  
  ]  
}
```

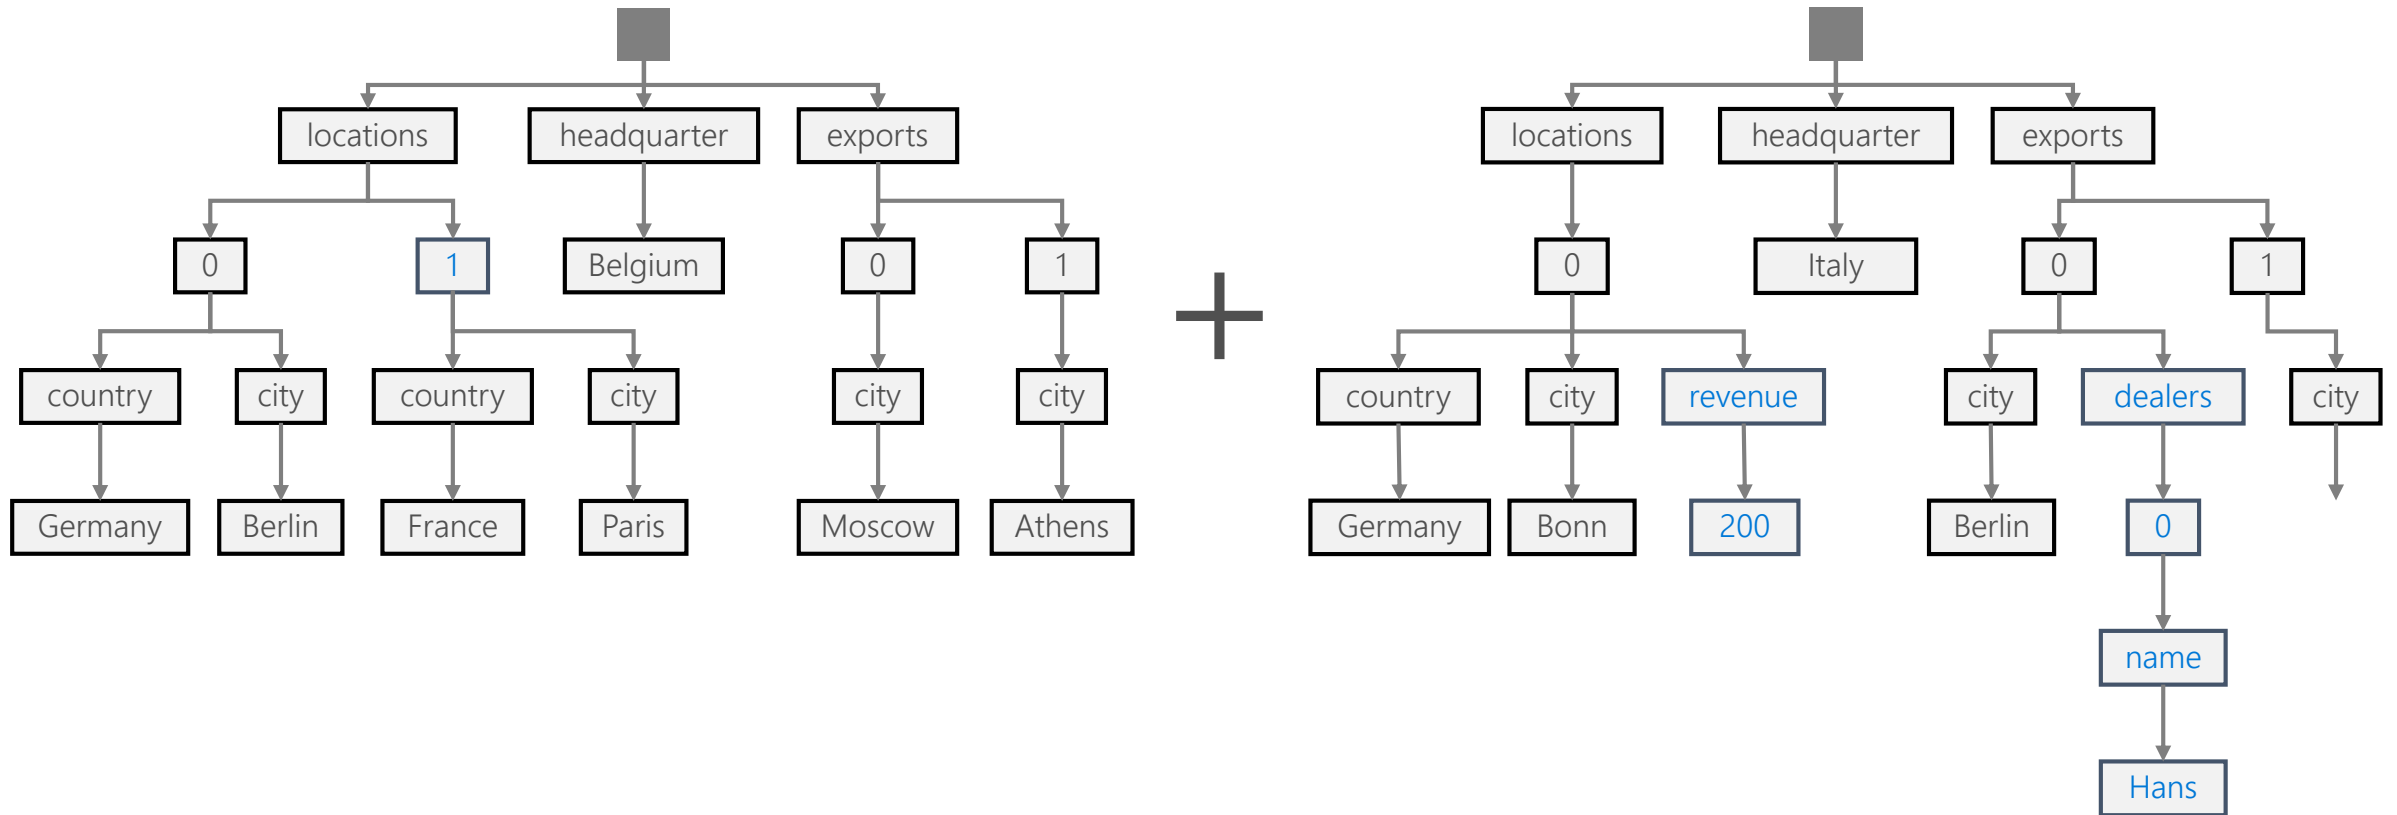


Indexing Json Documents - continued

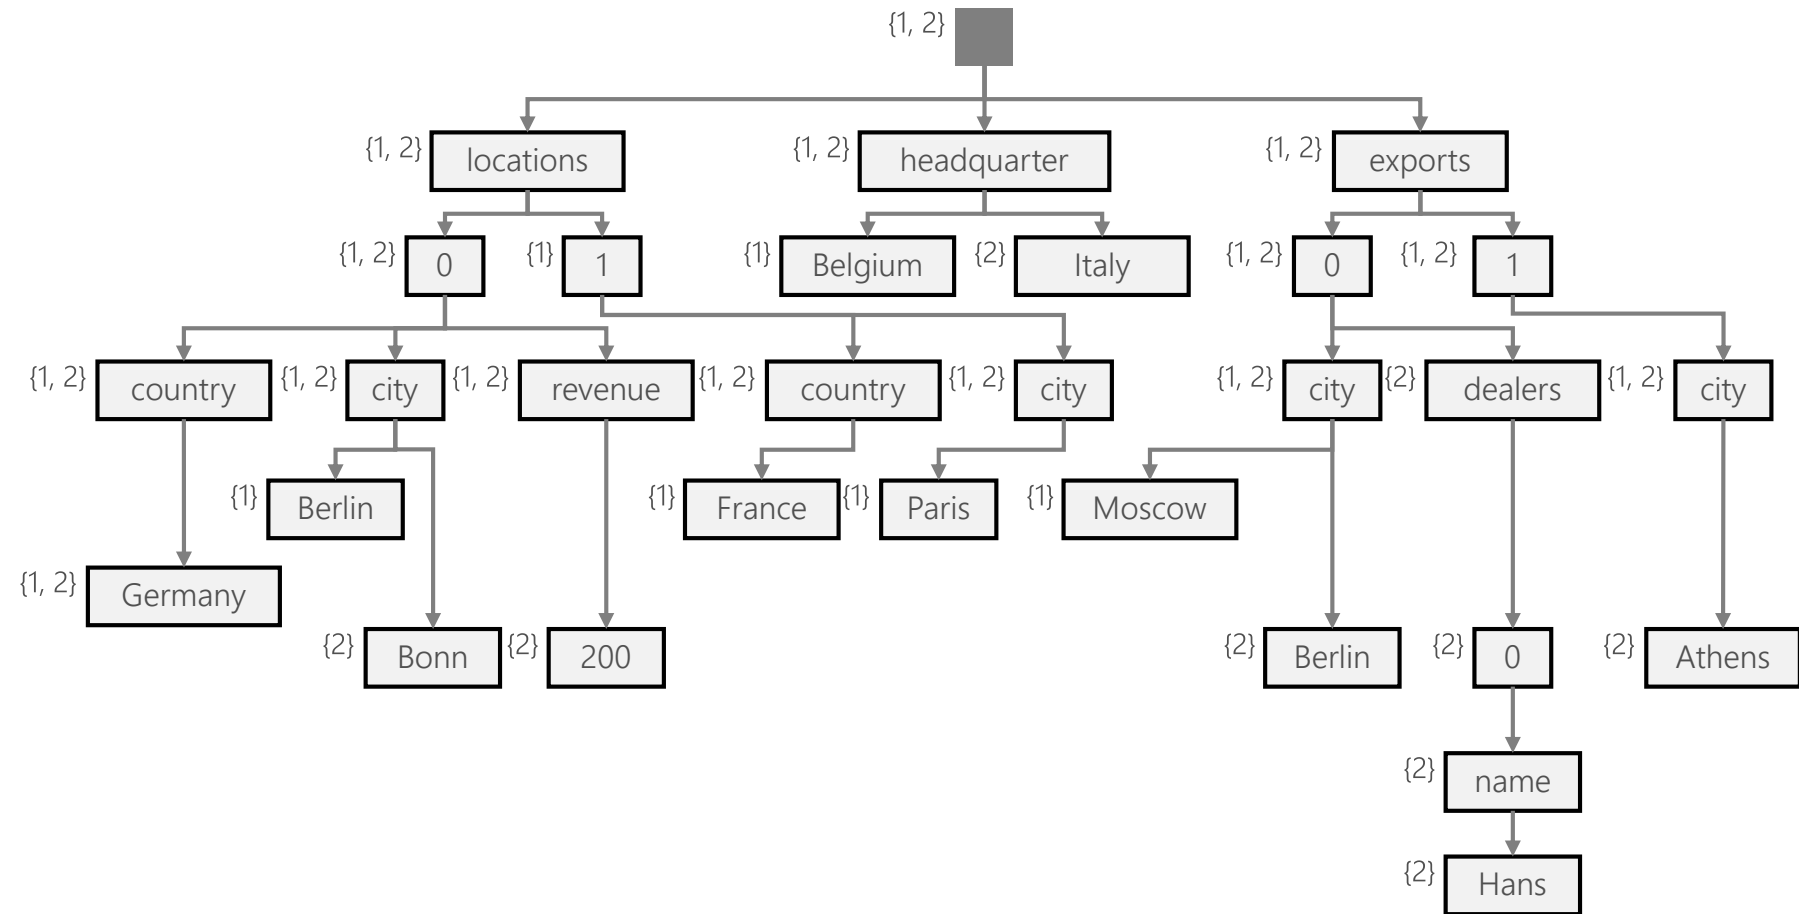
```
{  
  "locations": [  
    {  
      "country": "Germany",  
      "city": "Bonn",  
      "revenue": 200  
    }  
  ],  
  "headquarter": "Italy",  
  "exports": [  
    {  
      "city": "Berlin",  
      "dealers": [  
        { "name": "Hans" }  
      ]  
    },  
    { "city": "Athens" }  
  ]  
}
```



Indexing Json Documents - continued



Inverted Index



Indexing policies

CUSTOM INDEXING POLICIES

Though all Azure Cosmos DB data is indexed by default, you can specify a custom indexing policy for your collections.

Custom indexing policies allow you to design and customize the shape of your index while maintaining schema flexibility.

- Define trade-offs between storage, write and query performance, and query consistency
- Include or exclude documents and paths to and from the index
- Configure various index types

```
{
  "automatic": true,
  "indexingMode": "Consistent",
  "includedPaths": [{
    "path": "/*",
    "indexes": [{
      "kind": "Hash",
      "dataType": "String",
      "precision": -1
    }, {
      "kind": "Range",
      "dataType": "Number",
      "precision": -1
    }, {
      "kind": "Spatial",
      "dataType": "Point"
    }
  ]
}, {
  "excludedPaths": [{
    "path": "/nonIndexedContent/*"
  }]
}
```

Swiss Army Knife

Multi-model + Multi API
Reference Case



Multi-model + multi-API

- Different models:
 - Graph
 - Key-Value
 - Document DB
- API support:
 - SQL
 - JavaScript
 - Gremlin
 - MongoDB
 - Azure Table Storage
 - Cassandra



Mimicking Strategy



- MongoDB



- Cassandra



What model?

SQL API



Use the SQL API if you're building a new non-relational document database and want to query using familiar SQL syntax.

Gremlin API



Use the Gremlin API if you're building a graph database to model and traverse relationships among entities.

Table API



Use the Table API if you are migrating data from Azure Table storage to Azure Cosmos DB's premium table offering.

MongoDB API



Use the MongoDB API if you are migrating data from a MongoDB database to Azure Cosmos DB's fully managed service.

Cassandra API

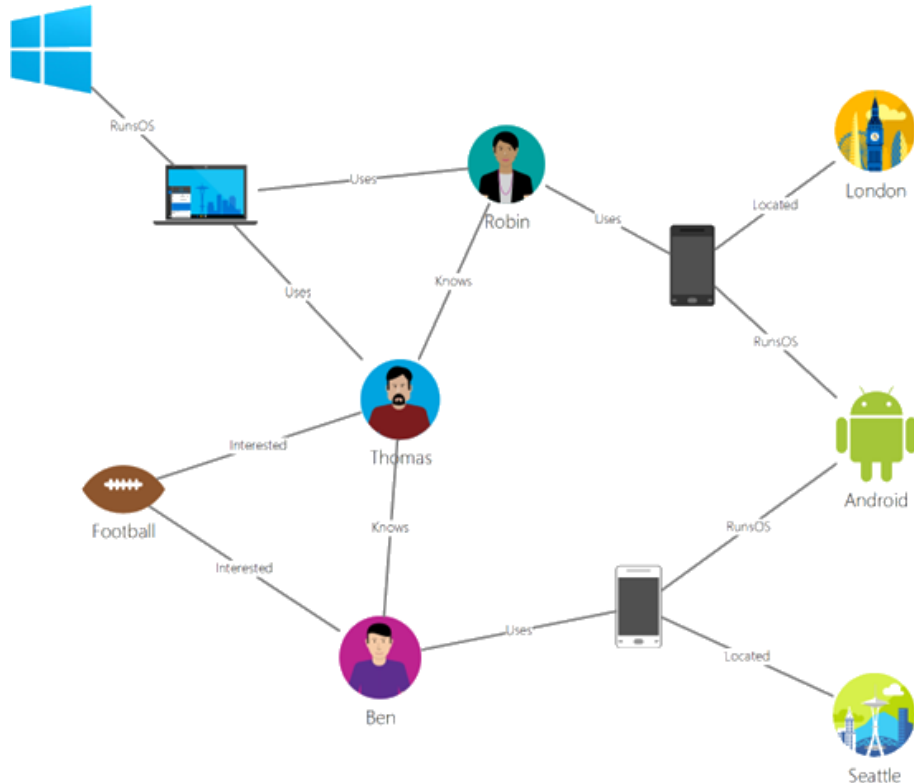


Use the Cassandra API if you are migrating data from Cassandra to Azure Cosmos DB's fully managed service.

Document (JSON)

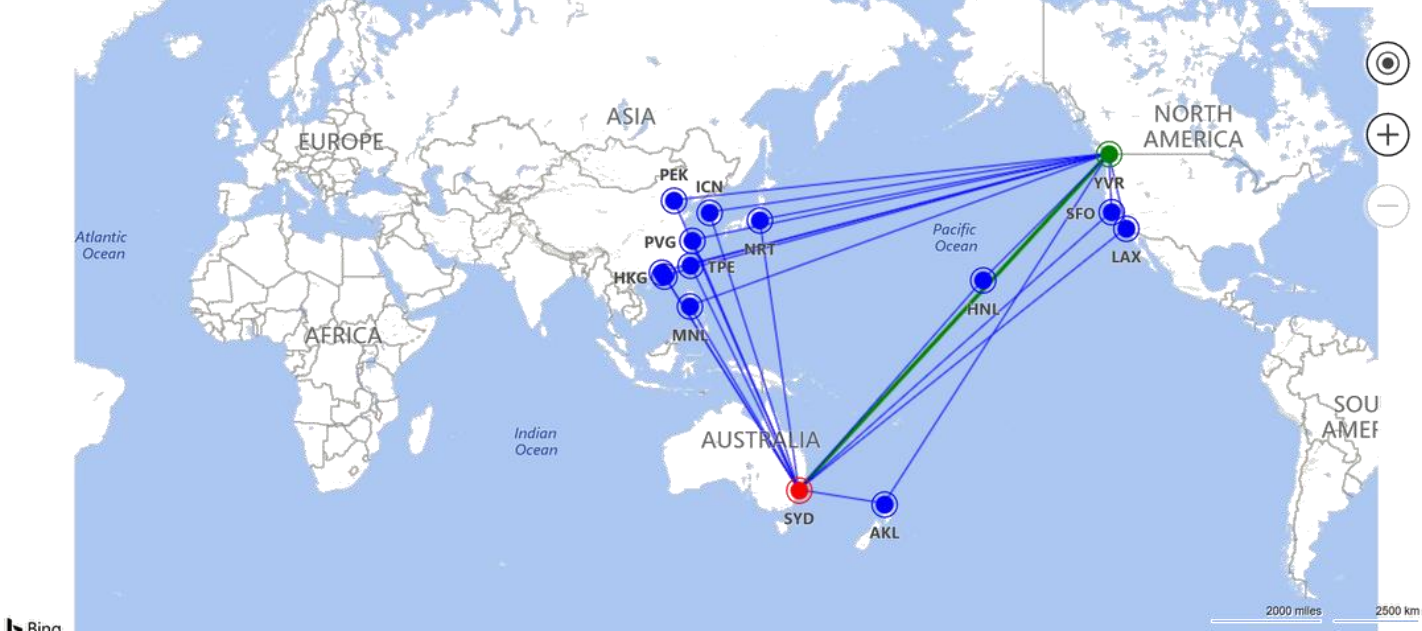
- A schema-less JSON database engine with rich SQL querying capabilities.
- Store documents
- Searchable by integrating with Azure Search
- Easy integration with Azure Functions
- Change Feed

Graph model



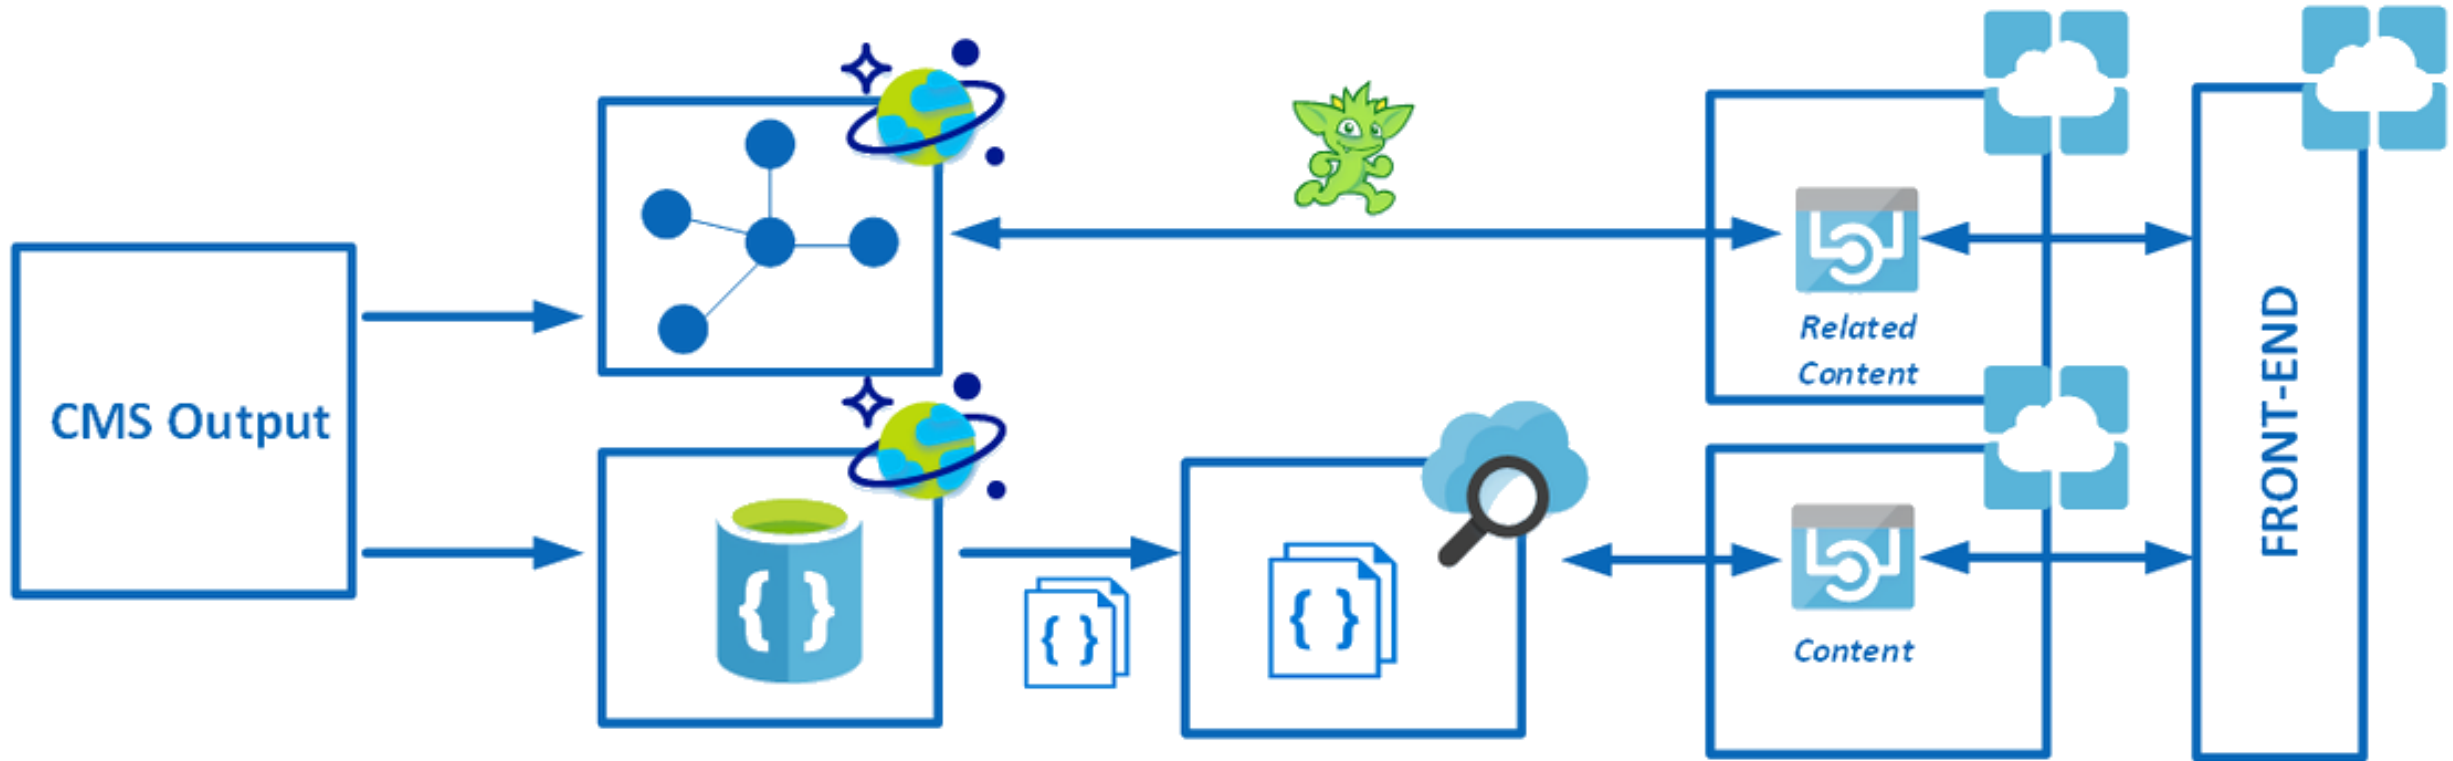
A **graph** is a structure that's composed of vertices and edges. Both **vertices** and **edges** can have an arbitrary number of properties.

- **Vertices** - Vertices denote discrete objects, such as a person, a place, or an event.
- **Edges** - Edges denote relationships between vertices. For example, a person might know another person, be involved in an event, and recently been at a location.
- **Properties** - Properties express information about the vertices and edges.



Demo - Graph

Use case – Cosmos DB Graph



<https://customers.microsoft.com/en-in/story/reed-business-information-professional-services-azure>

Key Takeaways

- Multiple options with models and API's
- Various consistency models
- Global scale
- Flexible through put
- Support for diverse scenario's



Call to action

- Documentation: <https://docs.microsoft.com/en-us/azure/>
- Middleware Friday: <https://www.youtube.com/watch?v=ZplZgOoGUzY>
- Graph demo: <https://github.com/anthonychu/cosmosdb-gremlin-flights>
- Pluralsight: <https://www.pluralsight.com/courses/azure-cosmos-db>
- Channel 9: <https://channel9.msdn.com/Events/Build/2018/BRK3319>

Build your skills with Microsoft Learn



Create an Azure Cosmos DB database built to scale



Insert and query data in your Azure Cosmos DB database



Build a .NET Core app for Azure Cosmos DB in Visual Studio Code

Running on time



Humana.
Rock'n'Roll
MARATHON & 1/2
ARIZONA
PHOENIX · SCOTTSDALE · TEMPE

JANUARY 19 - 20, 2019

OFFICIAL TIME

01:48:32



Thanks and
have a great
day!